

Seta Future — Phase-1 Kickoff Signoff


Field	Value
Document	Project signoff for kickoff
For	CEO · CTO · PMO
Author	Project Management Office
Kickoff (actual)	2026-04-23 (S1–S2 complete on provisional approval)
First demo	2026-05-20 (interim · end of S4)
MVP demonstration	2026-05-31 (staging only · 6-week window)

1. The Ask

1.1 Why this matters to SETA

- **Task & plan management** — plans, buckets, tasks, evidence-backed completion, 4 views, 4 personal hubs, bidirectional M365 Planner sync (detail in §5.1).
 - **AI assistant for every employee** — KB Q&A with citations, "my overdue tasks / why at-risk" answers, scheduled digests, approval-inbox-mediated writes, four-axis cost ceilings.
 - **Proves AI-leveraged velocity** — 2 non-trivial modules in 6 weeks with 6 FTE; S3 retro burn is the board-facing artefact for the rest of the SETA ecosystem build.
 - **Builds the platform core once** — kernel/RLS, identity, multi-tenancy, approval inbox, agent runtime, audit/replay, M365 connector, Terraform. Every later module (people · time · hiring · performance · projects · finance · goals) plugs in.
 - **SETA-internal pilot, staging only at MVP** — production cutover and external-tenant surface are post-MVP gates.
-

2. Architecture at a Glance



Invariants. Every table carries `tenant_id` with **RLS at the DB layer** · zones never query the DB directly (all data via tRPC) · **Terraform-only** infrastructure, no manual console changes · **ARM64 only** · secrets in AWS Secrets Manager.

3. Scope — What Ships by 2026-05-31

3.1 Committed deliverables (in scope)

Workstream	Headline
Agents module <i>(detail in §4)</i>	Conversational AI assistant — KB Q&A, role-scoped analysis, scheduled digests, approval-inbox-mediated writes, four-axis cost ceilings
Planner module <i>(detail in §5)</i>	Plans/buckets/tasks, evidence-backed completion, 4 view modes, 4 personal hubs, bidirectional Microsoft 365 Planner sync with conflict log

3.2 Acceptance criteria — demo on 2026-05-31

Live, end-to-end run on real SETA data **on the staging environment only** — no production deployment in Phase-1. Production cutover, agent write activation (Agents SAD §9.3), and full Planner launch (2026-06-08, Planner SAD §9.1) are post-MVP gates.

1. Magic-link or SSO login via Entra SSO.
2. **Plan & task management** — create plan + buckets + tasks (assignees · dates · labels · checklists · attachments); render Board, Grid, Charts, and Schedule views; My Day + My Tasks hubs load under caller scope.
3. **Mark task complete with evidence** → approval inbox → approve → propagates to M365 Planner via bidirectional sync; sync conflict surfaces in log with field-level diff.
4. **AI — KB Q&A** with handbook citations — ≥80% accuracy on a held-out rubric.
5. **AI — conversational Planner query & analysis** — natural-language questions on tasks, workload, and plan health, plus role-scoped throughput / balance / cross-team risk. Agent reads under caller scope, k-anonymity floor on team views, cites tasks.
6. **AI — inline copilot** in Planner with auto screen context; daily digest subscription delivers next business day.
7. **AI — conversational Planner control** — user issues task intents in natural language (create, update, reassign, reschedule, mark complete with evidence). Agent translates intent into a structured Planner write, routes through the approval inbox, executes on user approval under caller scope, and propagates to M365.
8. **Multi-tenant probe** — second tenant cannot see tenant 1's data; leak canary green.
9. **Failure modes** — cost-ceiling tripwire fires; KB miss returns plain-language refusal; sync retries idempotent.
10. Admin surfaces: sync-health · conflict log · cost ledger · audit query.

Gates: Board p95 ≤ 400 ms · My Tasks p95 ≤ 1.0 s · agent answer p95 ≤ 8 s · zero leak canary failures.

Quality bar: ≥70% unit-test coverage across lines, functions, and branches (CI-blocking — TDD per CLAUDE.md) · **main user flows covered by Playwright E2E (~20% of test surface)** — login + magic link, KB Q&A with citations, Planner task → evidence → approval → M365 sync, multi-tenant isolation probe.

Security gate (before staging hosts real data): secrets in AWS Secrets Manager · RLS policy-coverage test green · dep/container scan clean of criticals · M365 tokens encrypted with rotation · staff comms sent. Pen-test, DPIA, GDPR erasure → post-MVP.

3.3 Out of scope — programme-wide

Module-specific cuts in §4.3 and §5.2.

- **Other domain modules** — people, time, hiring, performance, projects, finance, goals
- **Manager / team-lead dashboards** and at-risk alerts

- **Meeting-transcript → plan conversion** (auto-generate plans/tasks from meeting scripts) — Phase-1.5 candidate
- **Multi-language UI · native mobile · multi-region · cross-tenant collab**
- **Slack / Teams chat surfaces · Outlook calendar**
- **External tenant onboarding** — SETA-internal pilot only

3.4 What 2026-05-31 is — and is not

✓ It is	✗ It is not
MVP demonstration on real SETA data	General Availability
Deployed to staging only	Deployed to production
PMO performance-evaluation anchor for the programme	External pilot launch
Foundation for the post-MVP Phase-1 GA gate	An all-modules platform
Internal-use ready (SETA staff, existing privacy)	External-tenant ready (GDPR erasure required)


4. Agents Module — Detailed Scope

Conversational AI assistant for every SETA employee. **MVP scope:** question-answering, role-scoped insights, and constrained writes (own-scope only, every write routed through the approval inbox). Autonomous writes are out.

The **approval inbox** is the safety mechanic: writes either preview inline (single-target) or route through the inbox (bulk, cross-target, destructive, or any write after free-text input). On approval, preconditions revalidate before execution; narrowed permissions fail the write with a structured event — never silently.

4.1 Runtime architecture — one turn, end to end

One turn = one Router plan up front, 1–2 sub-agents in parallel, one Synthesizer at the end. Every tool call traverses a single non-skippable **Tool Gateway** (permission · audit · cost). Every LLM call is recorded in the replay-keyed **Prompt Store** — any past turn reconstructable from its trace ID.



4.2 MVP capability buckets

Capability	What it does
Conversational surface	Global chat plus inline copilot inside Planner; per-user threads with streaming responses
Knowledge base Q&A	Ingest tenant handbooks / docs (Markdown · PDF · text); answer with citations; admin upload, edit, deprecate, re-index
Planner query & analysis	Natural-language questions on tasks, workload, plan health; role-scoped views (IC throughput · team-lead balance · org-leader cross-team risk) with k-anonymity on teams
Planner control (write)	Natural-language task intents → structured Planner write → routed through approval inbox → executed under caller scope on approval. Never silent, never autonomous at MVP
Scheduled digests	Opt-in morning brief · end-of-week status · stale-task nudges · at-risk alerts. Read-only or inbox-draft only
Cost & governance	Four-axis USD ceilings (per turn · user-day · tenant-day · delegation); admin controls model, ceilings, schedules, tool visibility
Audit & replay	Caller-identity audit on every action; deterministic replay from trace ID; plain-language refusal with reason (budget / unavailable / refused)

Internals (Router / sub-agents / Synthesizer · layered memory · ReAct loops · circuit breakers) are specified in the Agents SAD; FR-level refs in the SDLC traceability matrix.

4.3 Agents-specific deferrals

- **Cross-conversation memory** — within-conversation context only (no "you said X yesterday")
 - **Autonomous writes** — scheduled / event-triggered runs cannot write without inbox approval
 - **Event-triggered runs** — cron only at MVP; outbox-event subscriptions pre-designed, not wired
 - **Meeting-transcript → plan conversion** — auto-generation of plans/tasks from meeting scripts is Phase-1.5
 - **Image / OCR ingestion** — text KB only
 - **Multi-provider routing** — single-provider (OpenAI); abstraction is post-Phase-1
 - **LLM-as-judge gating** — deferred until $\geq 95\%$ inter-rater corpus; deterministic scorers only
-

5. Planner Module — Detailed Scope

A work-tracking surface that **complements** Microsoft 365 Planner today and is positioned to **replace it** for tenants without an existing M365 investment — Future adds bidirectional sync, evidence-backed completion, personal hubs, and audit history that M365 Planner does not provide. Additional integrations (Trello, Asana, Jira) are post-MVP; the architecture is provider-agnostic by design.

5.1 MVP capability buckets

Capability	What it does
Plans, buckets, tasks	Team plans (multi-member) and personal plans (auto-provisioned, private); ordered buckets; tasks with assignees, dates, labels, checklists, attachments, comments; soft-delete + audit
Four view modes	Board (kanban) · Grid (sortable table) · Charts (by bucket / assignee / progress / priority) · Schedule (timeline by day / week / month). Single filter+search across all four. Not a Gantt
Four personal hubs	My Day · My Tasks · Personal Charts · Carry-Over (auto roll-forward of unfinished pins)
Evidence model	First-class records (file · link · note); states unsubmitted / submitted / verified / rejected; tracked independently of task completion
M365 bidirectional sync	Tenant OAuth · adaptive pull + idempotent push · last-write-wins conflict log with both before-states · dry-preview on first sync · credential revoke pauses sync within one cycle
Admin surface	M365 connect / disconnect · conflict log with field-level diff and per-task force-resync · "what changed and why" query · daily sync-health summary · attachment quotas

Field-level limits, FR-PL refs, plan-type fixing, and Graph delta-query specifics are in the Planner SAD.


5.2 Planner-specific deferrals

- **M365 Planner Premium-tier** — custom fields, conditional colouring, People view, sprints/backlog, custom calendars, M365 Copilot, rich-text
- **Gantt** — no dependencies, critical path, or milestones (Schedule = timeline-by-date)
- **Real-time presence / cursors** — async only
- **AI reminders / EOD digests** — live in **Agents**, not Planner
- **Attachment processing** — no OCR, virus scan, or redaction
- **Per-field merge** — last-write-wins fixed
- **Copy/template · Excel/CSV export · per-bucket colour** — Phase-1.5 candidates
- **Push channels** — pull-only; no WebSocket / SSE
- **Other task systems** — M365 only (no Asana / Trello / Jira)

5.3 Hard external dependencies

Dependency	Owner	Deadline
M365 sandbox tenant + Entra admin consent + Graph scopes	SETA IT	2026-05-08
Secrets manager with online rotation for M365 credentials	Platform Security	S3
Identity directory sync (users, groups)	Identity / People	S3 D1 (05-07)

6. Timeline — Six One-Week Sprints



Sprint	Window	Goal	Gate
S1	04-23 → 04-29	Monorepo + toolchain — Done	—
S2	04-30 → 05-06	Backend, frontend, auth, DB skeletons + SSO — Done	—
S3	05-07 → 05-13	Staging deployable · People + Planner CRUD + Agents chat skeleton · 4 cross-module API contracts locked Day 1	Velocity re-baseline retro
S4	05-14 → 05-20	Planner views + hubs + sync core · Agents KB + retrieval-augmented Q&A + write/execute mode · staging hardened	First demo — 2026-05-20
S5	05-21 → 05-27	MVP feature-complete + linking · code freeze at S5 close	Go / no-go gate for S6
S6	05-28 → 05-31 (4 days)	Hardening only — bug fixes, performance / accessibility, security, demo prep, requirements-traceability check	MVP demo — 2026-05-31

6.1 Load-bearing dates

If any slips, the demo slips.

- **S3 D2 (05-08)**. M365 sandbox + Entra consent + Graph scopes (SETA IT); KB ingestion pipeline complete (Agents track).
- **S3 retro (05-13)**. Velocity re-baseline — target ~45 SP/eng/sprint; below plan → §6.2 cut order applies.
- **End of S4 (05-20)**. **First demo** — interim walkthrough on staging; Planner views + hubs + KB Q&A working end-to-end. Surfaces gaps before code freeze.
- **End of S5 (05-27)**. Code freeze. Linking complete; integration smoke must be green or S6 hardening cannot land the demo.

6.2 Phase-1.5 cut order (pre-agreed, applies at S3 retro)

Apply in order until scope fits velocity: (1) scheduled digests · (2) bidirectional sync → one-way push · (3) Charts + Schedule views · (4) Personal Charts hub · (5) agent writes (already gated in Agents SAD §9.3).

6.3 Why one-week sprints

Six checkpoints, not three. AI-leveraged velocity (~45 SP/eng/sprint) is a forecast — S3 retro re-baselines from real burn. Two-week sprints would only surface the miss at week 5; too late to cut.

7. Team — 6.0 FTE for 6 weeks

Role	Allocation	Status	Owns
Planner Module	1 fullstack	Existing 0.5 + need to hire 0.5	Plans/tasks/evidence → views/hubs → M365 sync → linking · web-admin Planner surfaces · People exact-subject resolver
Agents Module	1 AI + 1 fullstack	Existing 1 AI · need to hire 1 fullstack	Chat → KB + writes via approval inbox → admin/governance → linking · web-admin KB surfaces
Deployment	1 DevOps	Existing (needs IT support)	Staging environment · web-shell SSO + magic-link · dual-tenant probe
QA Engineer	1.0 FTE	To hire — onboard immediately	Manual + exploratory · regression · launch-gate verification
Business Analyst	0.5 FTE	Existing	Requirements · acceptance · traceability · pilot feedback · doc work-stream owner
Scrum Master	0.5 FTE	To hire — onboard immediately	Cadence · retros · impediments · velocity
Total	6.0 FTE for 6 weeks	To hire — 1.5 fullstack + 1 QA + 0.5 SM	

Bus factor. Planner = single fullstack; Agents = single AI engineer. Either absence > 3 days → §6.2 cut activates (see §8 R3, R4).

Enabling support (not against capacity): IT/DevOps provides AWS, DNS, certs, M365 sandbox + Entra consent, shared Terraform + IAM. All pre-kickoff.

8. Risks That Bear on Scope or Timeline

The six below need a CEO/CTO/PMO ruling **before** signoff. Full register in the board paper.

#	Risk	P/I	Mitigation / Action
R1	AI-leveraged velocity (~45 SP/eng/sprint) is a forecast, not measured	M / H	S3 retro re-baselines from real burn; below forecast → §6.2 cut order
R2	Hires not yet onboard (1.5 FS + 1 QA + 0.5 SM) — already late vs. 05-07 target	H / H	Onboard immediately; each week of FS slip ≈ one §6.2 cut
R3	Single-fullstack Planner — bus factor	M / H	PR reviewer named; absence > 3 days → §6.2 cut #2
R4	Single-AI-engineer Agents — bus factor	M / H	FS cross-trained on chat + KB in S2; AI absence > 3 days → reduce Agents to KB Q&A + chat (drop Planner analysis + writes)
R5	M365 Graph throttling untested at SETA scale	M / M	S4 backfill rehearsal (Planner SAD §9.2); adaptive cadence + pre-flight limits
R6	S5 overload — feature-complete + linking same week	M / H	Code-freeze threshold mid-week; if behind, late-cut Planner UI polish + Agents model-degradation ladder (per SDLC plan risk #9)

9. Budget — Cloud + AI Envelope

Line item	USD/month	Status
AWS infrastructure — staging during MVP (\$127); prod (\$349) activates at post-MVP cutover	~\$476 (GA)	Budgeted
Claude Code Max subscriptions (developer tooling)	~\$220	Budgeted
OpenAI inference — SETA-scale forecast (300 staff · ~10 turns/user/day · 4-axis cost ceilings active)	~\$880	Budgeted
Team payroll (6.0 FTE × 6 weeks)	(existing)	Already approved
Recommended monthly operating envelope (cap)	~\$2,100	Approve

One-offs: M365 sandbox (zero marginal); KB ingestion of SETA corpus (~\$0.20).

Variance trigger: per-tenant ≥ 2× ceiling on rolling 7-day → budget review; ≥ 3× → admin alert + pricing review.

10. Decisions Required — Sign by 2026-05-13 (S3 retro)

S1–S2 were executed on provisional approval. Signing now **ratifies that work and authorises S3–S6**. Latest defensible signoff is the **S3 retro on 2026-05-13** — beyond that, the §6.2 cut order applies without re-escalation.

By signing, approvers confirm each of:

1. **Ratify** the 2026-04-23 commencement and the S1–S2 work delivered on provisional approval.
2. **Authorise S3–S6** with the scope, timeline, team, and budget above.
3. **Authorise outstanding hires immediately** — 1.5 fullstack + 1 QA + 0.5 SM. **Any further hiring delay slips the timeline** (already late vs. original $\leq$ 2026-05-07 target — see R2). Each week of slip on FS hire $\approx$ one cut from §6.2.
4. **Confirm M365 dependency** — sandbox + Entra consent + Graph scopes due **today, 2026-05-08** (S3 D2). Slip $\rightarrow$ §6.2 cut #2 (sync $\rightarrow$ one-way push) activates.
5. **Confirm work-stream fold-ins** per §7 (People · web-admin · web-shell · Documentation). 6.0 FTE assumes these; reject $\rightarrow$ scope re-baselines.
6. **Ring-fence the team** for the remaining four sprints — no client-services pulls. AI-velocity bet is void otherwise.
7. **Pre-agree the §6.2 cut order** — applies at S3 retro without further escalation.
8. **Acknowledge agile delivery** — scope may be refined mid-sprint for clarity, quality, or improvement opportunities surfaced during implementation. **Material changes (in/out of §3.1, demo-affecting) are communicated to CEO/CTO/PMO within 24h** with rationale and impact; non-material refinements land in the sprint demo.
9. **Acknowledge** 2026-05-31 = MVP demo on **staging**, not GA. Production cutover, agent writes (Agents SAD §9.3), and Planner launch (2026-06-08, Planner SAD §9.1) are post-MVP gates.

11. Sign-off

Role	Name	Signature	Date
CEO	Hung Vu		YYYY-MM-DD
CTO	Thu Mai		YYYY-MM-DD
PMO Lead	Hoang Nguyen		YYYY-MM-DD